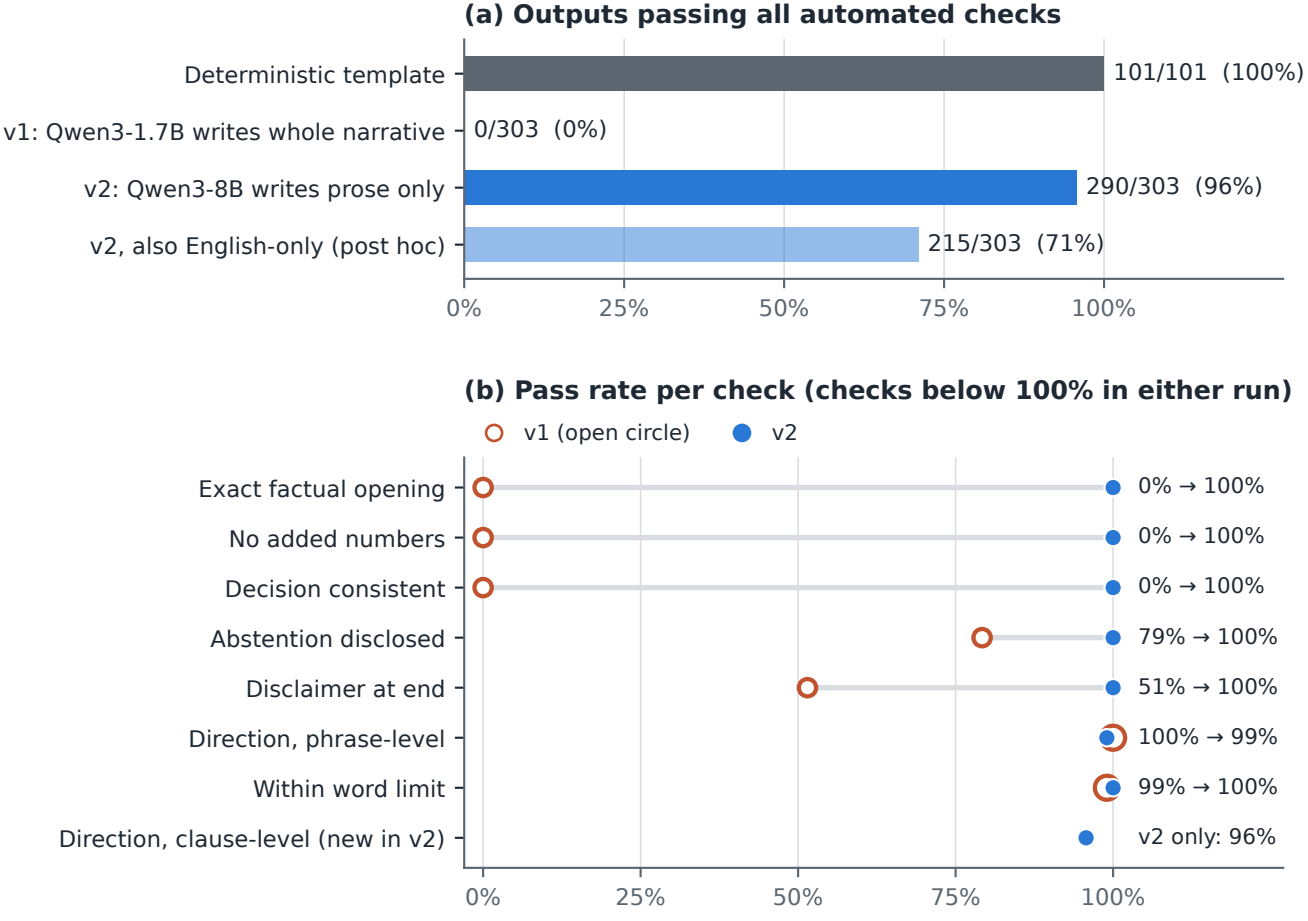
(a) Outputs passing all automated checks
Deterministic template
101/101 (100%)
v1: Qwen3-1.7B writes whole narrative
0/303 (0%)
v2: Qwen3-8B writes prose only
290/303 (96%)
v2, also English-only (post hoc)
215/303 (71%)
0%
25%
50%
75%
100%
(b) Pass rate per check (checks below 100% in either run)
v1 (open circle)
v2
Exact factual opening
0% → 100%
No added numbers
0% → 100%
Decision consistent
0% → 100%
Abstention disclosed
79% → 100%
Disclaimer at end
51% → 100%
Direction, phrase-level
100% → 99%
Within word limit
99% → 100%
Direction, clause-level (new in v2)
v2 only: 96%
0%
25%
50%
75%
100%